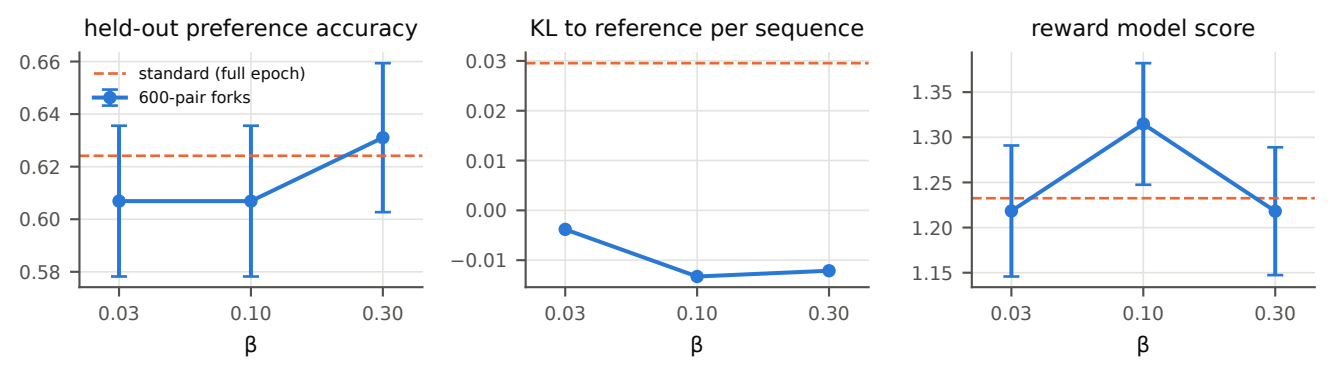
held-out preference accuracy
standard (full epoch)
600-pair forks
0.66
0.64
0.62
0.60
0.58
0.03
0.10
0.30
β
KL to reference per sequence
0.03
0.02
0.01
0.00
−0.01
0.03
0.10
0.30
β
reward model score
1.35
1.30
1.25
1.20
1.15
0.03
0.10
0.30
β