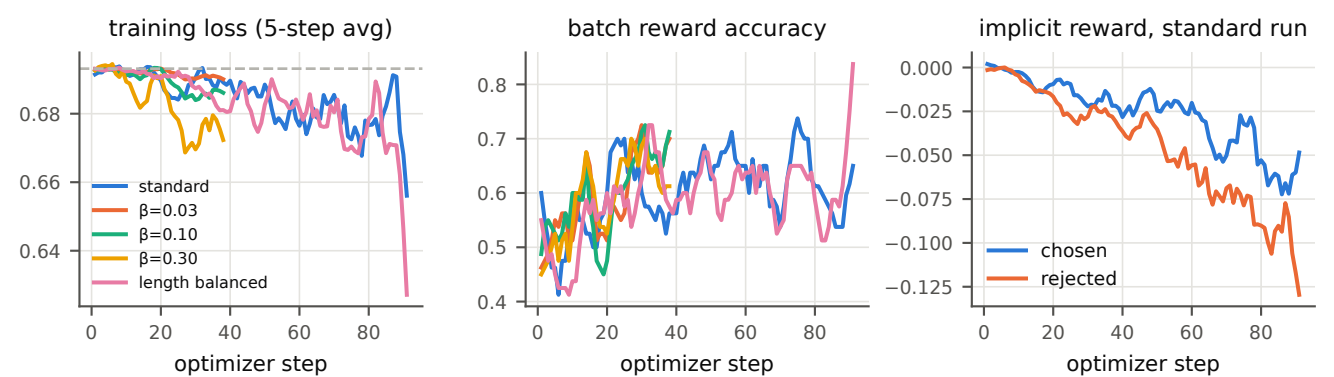
training loss (5-step avg)
0.68
0.66
0.64
standard
β=0.03
β=0.10
β=0.30
length balanced
0
20
40
60
80
optimizer step
batch reward accuracy
0.8
0.7
0.6
0.5
0.4
0
20
40
60
80
optimizer step
implicit reward, standard run
0.000
−0.025
−0.050
−0.075
−0.100
−0.125
chosen
rejected
0
20
40
60
80
optimizer step