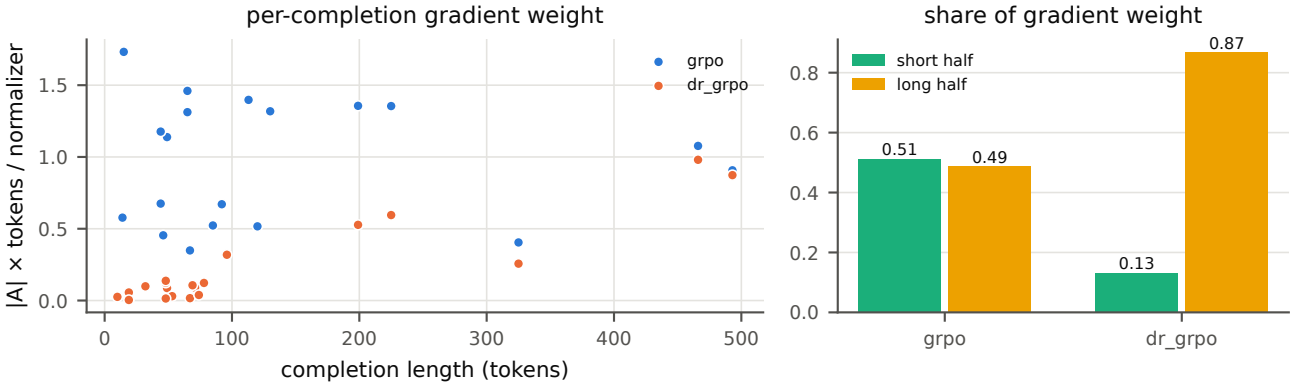
per-completion gradient weight
|A| × tokens / normalizer
completion length (tokens)
grpo
dr_grpo
0.0
0.5
1.0
1.5
0
100
200
300
400
500
share of gradient weight
short half
long half
0.51
0.49
0.13
0.87
0.0
0.2
0.4
0.6
0.8
grpo
dr_grpo